

In [20]:

```
import numpy as np
import pymc as pm
import matplotlib.pyplot as plt
```

In [22]:

```
# Generate synthetic data
np.random.seed(42)
X = np.linspace(0, 10, 50)
true_intercept = 1
true_slope = 2
Y = true_intercept + true_slope * X + np.random.normal(scale=1, size=X.size)
```

In [24]:

```
# Bayesian linear regression model
with pm.Model() as model:
    # Priors for unknown model parameters
    intercept = pm.Normal("Intercept", mu=0, sigma=10)
    slope = pm.Normal("Slope", mu=0, sigma=10)
    sigma = pm.HalfNormal("Sigma", sigma=1)

    # Expected value of outcome
    Y_pred = intercept + slope * X

    # Likelihood of observations
    Y_obs = pm.Normal("Y_obs", mu=Y_pred, sigma=sigma, observed=Y)

    # Sampling posterior
    trace = pm.sample(2000, return_inferencedata=True, progressbar=True)
```

Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt\_diag...

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [Intercept, Slope, Sigma]

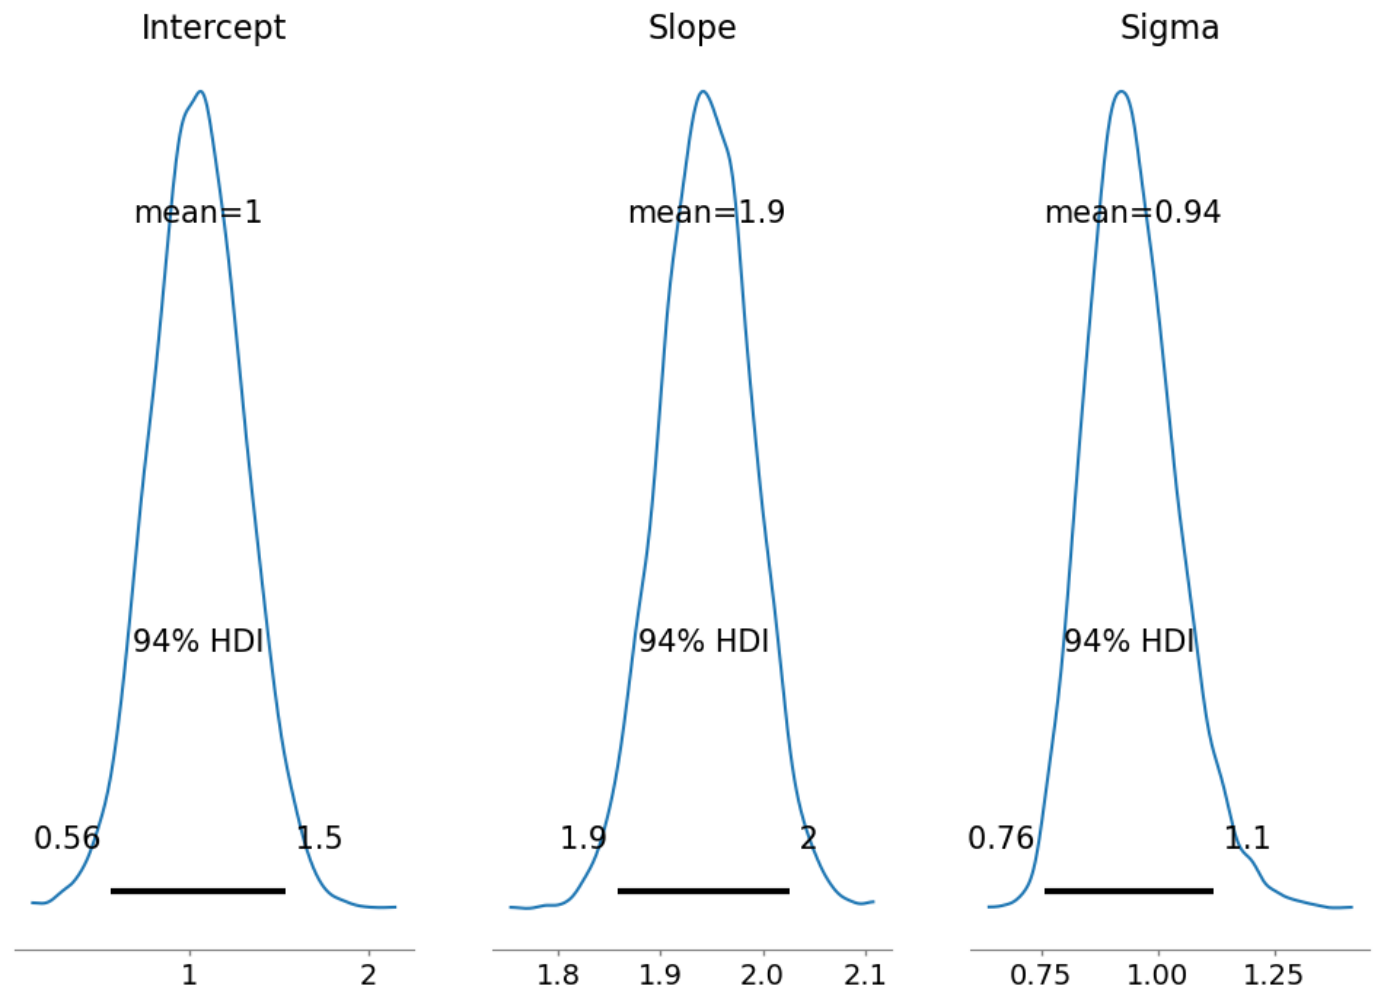
Output()

Sampling 4 chains for 1\_000 tune and 2\_000 draw iterations (4\_000 + 8\_000 draws total) took 31 seconds.

In [25]:

```
# Plot posterior distributions
pm.plot_posterior(trace, figsize=(12, 8))
plt.show()

# Summary of the posterior
print(pm.summary(trace, hdi_prob=0.95))
```



	mean	sd	hdi_2.5%	hdi_97.5%	mcse_mean	mcse_sd	ess_bulk	\
Intercept	1.049	0.261	0.558	1.575	0.004	0.003	3622.0	
Slope	1.944	0.045	1.852	2.029	0.001	0.001	3662.0	
Sigma	0.945	0.098	0.760	1.136	0.001	0.001	4884.0	

	ess_tail	r_hat
Intercept	3361.0	1.0
Slope	3375.0	1.0
Sigma	4498.0	1.0